

Linear Regression by Optimisation

leta199

May 28, 2026

Introduction

Linear models act to describe the relationship between a response variable Y and its predictor X . It acts as the best linear approximation to the true relationship.

The true model is represented by the equation:

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon$$

where:

- Y_i – i -th response variable
- X_i – i -th predictor variable
- β_0 – true intercept term
- β_1 – true predictor coefficient
- ϵ – random error

Using collected data, we can make the best point estimates for the true value of coefficients and have:

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i + e_i$$

where:

- $\hat{\beta}_0$ – estimated intercept term
- $\hat{\beta}_1$ – estimated predictor coefficient
- e_i – residual

Expected Value of Y

The expected value of the response variable Y given a specific predictor value X_i represents the population regression line. Because the expected value of the random error term is assumed to be zero ($E[\epsilon] = 0$), the expected value of our estimate reduces purely to the deterministic model component:

$$E[Y_i] = \hat{\beta}_0 + \hat{\beta}_1 X_i$$

This equation provides the predicted mean value of Y for any given value of X , effectively removing the individual sample noise (e_i or ϵ) from the evaluation.

Using Gradient Descent

In order to find the best approximations $\hat{\beta}_0$ & $\hat{\beta}_1$, we can use gradient descent and minimise mean aggregate errors known as mean squared error (MSE). MSE is given by the formula:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

We can utilise partial derivatives to find estimates of β_0 & β_1 . By expanding $\hat{Y}_i$ as our prediction function, our objective cost function $f(\hat{\beta}_0, \hat{\beta}_1)$ becomes:

$$f(\hat{\beta}_0, \hat{\beta}_1) = \frac{1}{n} \sum_{i=1}^n \left[Y_i - (\hat{\beta}_0 + \hat{\beta}_1 X_i) \right]^2$$

Taking the partial derivatives with respect to each parameter yields:

$$\frac{\partial f}{\partial \hat{\beta}_0} = -\frac{2}{n} \sum_{i=1}^n \left[Y_i - (\hat{\beta}_0 + \hat{\beta}_1 X_i) \right] \quad \frac{\partial f}{\partial \hat{\beta}_1} = -\frac{2}{n} \sum_{i=1}^n X_i \left[Y_i - (\hat{\beta}_0 + \hat{\beta}_1 X_i) \right]$$

Now to reach optimal values of $\hat{\beta}_0$ & $\hat{\beta}_1$ that minimise errors e_i , where $e_i = Y_i - \hat{Y}_i$, we use gradient descent iteratively by taking steps L in the negative direction of the gradient.

$$\hat{\beta}_{0(\text{new})} = \hat{\beta}_0 - L \frac{\partial f}{\partial \hat{\beta}_0} \quad \hat{\beta}_{1(\text{new})} = \hat{\beta}_1 - L \frac{\partial f}{\partial \hat{\beta}_1}$$